

Mudcard

- **Why do we need a validation set if we already have a test set?**
 - the validation set is used for model selection
 - the test set is used to estimate how well the model is expected to perform on previously unseen data points once it is deployed
- **Is there a best practice for reading the Lasso and Ridge regressions directly from the graph?**
 - I'm not sure what you are referring to.
 - Please expand on it and submit your question on Ed
- **Wondering why do we want coeffs to be exactly zero for feature selection?**
 - If a coefficient is 0, the corresponding feature does not contribute to the prediction of the linear model
 - $y' = \langle w, X \rangle = \sum (w_1 * X_1 + w_2 * X_2 + w_3 * X_3 + \dots)$
 - If w_i is 0, the corresponding feature X_i is not used in the prediction because $w_i * X_i = 0$ if $w_i = 0$
- **How can we interpret the learned coefficients from LogisticRegression?**
 - See PS5 problem 1 for the answer

Lecture 9: CV and Intro to Interpretability

By the end of this lecture, you will be able to

- perform basic hyperparameter tuning
- apply GridSearchCV
- describe why interpretability is important

The supervised ML pipeline

0. Data collection/manipulation: you might have multiple data sources and/or you might have more data than you need

- you need to be able to read in datasets from various sources (like csv, excel, SQL, parquet, etc)
- you need to be able to filter the columns/rows you need for your ML model
- you need to be able to combine the datasets into one dataframe

1. Exploratory Data Analysis (EDA): you need to understand your data and verify that it doesn't contain errors

- do as much EDA as you can!

2. Split the data into different sets: most often the sets are train, validation, and test (or holdout)

- practitioners often make errors in this step!
- you can split the data randomly, based on groups, based on time, or any other non-standard way if necessary to answer your ML question

3. Preprocess the data: ML models only work if X and Y are numbers! Some ML models additionally require each feature to have 0 mean and 1 standard deviation (standardized features)

- often the original features you get contain strings (for example a gender feature would contain 'male', 'female', 'non-binary', 'unknown') which needs to be transformed into numbers
- often the features are not standardized (e.g., age is between 0 and 100) but it needs to be standardized

4. Choose an evaluation metric: depends on the priorities of the stakeholders

- often requires quite a bit of thinking and ethical considerations

5. Choose one or more ML techniques: it is highly recommended that you try multiple models

- start with simple models like linear or logistic regression
- try also more complex models like nearest neighbors, support vector machines, random forest, etc.

6. Tune the hyperparameters of your ML models (aka cross-validation or hyperparameter tuning)

- ML techniques have hyperparameters that you need to optimize to achieve best performance
- for each ML model, decide which parameters to tune and what values to try
- loop through each parameter combination
 - train one model for each parameter combination
 - evaluate how well the model performs on the validation set
- take the parameter combo that gives the best validation score
- evaluate that model on the test set to report how well the model is expected to perform on previously unseen data

7. Interpret your model: black boxes are often not useful

- check if your model uses features that make sense (excellent tool for debugging)

- often model predictions are not enough, you need to be able to explain how the model arrived to a particular prediction (e.g., in health care)

By the end of this lecture, you will be able to

- **perform basic hyperparameter tuning**
- apply GridSearchCV
- describe why interpretability is important

Let's put everything together!

import packages

load your dataset

create feature matrix and target variable

for i in random_states:

- split the data
- preprocess it
- decide which hyperparameters you'll tune and what values you'll try
- for combo in hyperparameters:
 - train your ML algo
 - calculate training scores
 - calculate validation scores
- select best model based on the mean and std validation scores
- predict the test set using the best model
- return your test score (generalization error)
- return the best model

The deliverables of an ML pipeline are:

- at the very least:
 - the n number of best models you selected based on the validation scores
 - yes, you need ALL of them
 - the n number of test scores (or you can calculate mean and stdev test scores)
- in some cases, you need your whole pipeline in a reproducible format like a container
 - others might want to see how you preprocess and split your data
 - what models you try
 - your evaluation metric
 - they might want to see your train and validation scores to verify the range of your hyperparameters

- etc.

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder, OrdinalEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import ParameterGrid
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split

df = pd.read_csv('../data/adult_data.csv')

# let's separate the feature matrix X, and target variable y
y = df['gross-income'] # remember, we want to predict who earns more than 50k
X = df.loc[:, df.columns != 'gross-income'] # all other columns are features

# collect which encoder to use on each feature
# needs to be done manually
ordinal_ftrs = ['education']
ordinal_cats = [[' Preschool', ' 1st-4th', ' 5th-6th', ' 7th-8th', ' 9th', ' 10th',
                ' Some-college', ' Assoc-voc', ' Assoc-acdm', ' Bachelors', ' Masters',
                ' Doctorate'],
                ['workclass', 'marital-status', 'occupation', 'relationship', 'race']]
onehot_ftrs = ['workclass', 'marital-status', 'occupation', 'relationship', 'race']
minmax_ftrs = ['age', 'hours-per-week']
std_ftrs = ['capital-gain', 'capital-loss']

# collect all the encoders into one preprocessor
preprocessor = ColumnTransformer(
    transformers=[
        ('ord', OrdinalEncoder(categories = ordinal_cats), ordinal_ftrs),
        ('onehot', OneHotEncoder(sparse_output=False, handle_unknown='ignore'), onehot_ftrs),
        ('minmax', MinMaxScaler(), minmax_ftrs),
        ('std', StandardScaler(), std_ftrs)])

prep = Pipeline(steps=[('preprocessor', preprocessor)]) # for now we only pr
```

Quiz

Let's recap preprocessing. Which of these statements are true?

Basic hyperparameter tuning

```
In [2]: help(ParameterGrid)
```

Help on class ParameterGrid in module sklearn.model_selection._search:

```
class ParameterGrid(builtins.object)
|   ParameterGrid(param_grid)
|
|   Grid of parameters with a discrete number of values for each.
|
|   Can be used to iterate over parameter value combinations with the
|   Python built-in function iter.
|   The order of the generated parameter combinations is deterministic.
|
|   Read more in the :ref:`User Guide <grid_search>`.
|
|   Parameters
|   -----
|   param_grid : dict of str to sequence, or sequence of such
|       The parameter grid to explore, as a dictionary mapping estimator
|       parameters to sequences of allowed values.
|
|       An empty dict signifies default parameters.
|
|       A sequence of dicts signifies a sequence of grids to search, and is
|       useful to avoid exploring parameter combinations that make no sense
|       or have no effect. See the examples below.
|
|   Examples
|   -----
|   >>> from sklearn.model_selection import ParameterGrid
|   >>> param_grid = {'a': [1, 2], 'b': [True, False]}
|   >>> list(ParameterGrid(param_grid)) == (
|   ...     [{'a': 1, 'b': True}, {'a': 1, 'b': False},
|   ...     {'a': 2, 'b': True}, {'a': 2, 'b': False}])
|   True
|
|   >>> grid = [{'kernel': ['linear']}, {'kernel': ['rbf'], 'gamma': [1, 1
0]}]
|   >>> list(ParameterGrid(grid)) == [{'kernel': 'linear'},
|   ...                               {'kernel': 'rbf', 'gamma': 1},
|   ...                               {'kernel': 'rbf', 'gamma': 10}]
|   True
|   >>> ParameterGrid(grid)[1] == {'kernel': 'rbf', 'gamma': 1}
|   True
|
|   See Also
|   -----
|   GridSearchCV : Uses :class:`ParameterGrid` to perform a full parallelize
d
|       parameter search.
|
|   Methods defined here:
|
|   __getitem__(self, ind)
|       Get the parameters that would be ``ind``th in iteration
|
|   Parameters
|   -----
```

```

    ind : int
        The iteration index

    Returns
    -----
    params : dict of str to any
        Equal to list(self)[ind]

__init__(self, param_grid)
    Initialize self. See help(type(self)) for accurate signature.

__iter__(self)
    Iterate over the points in the grid.

    Returns
    -----
    params : iterator over dict of str to any
        Yields dictionaries mapping each estimator parameter to one of i
ts
        allowed values.

__len__(self)
    Number of points on the grid.

-----
Data descriptors defined here:

__dict__
    dictionary for instance variables

__weakref__
    list of weak references to the object

```

```

In [3]: # let's train a random forest classifier

# we will loop through nr_states random states so we will return nr_states t
nr_states = 5
test_scores = np.zeros(nr_states)
final_models = []

# loop through the different random states
for i in range(nr_states):
    print('randoms state '+str(i+1))

    # first split to separate out the training set
    X_train, X_other, y_train, y_other = train_test_split(X,y,train_size = 0

    # second split to separate out the validation and test sets
    X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_si

    # preprocess the sets
    X_train_prep = prep.fit_transform(X_train)
    X_val_prep = prep.transform(X_val)
    X_test_prep = prep.transform(X_test)

```

```

# decide which parameters to tune and what values to try
# the default value of any parameter not specified here will be used
param_grid = {
    'penalty': ['l1'],
    'C': np.logspace(-3,3,7), # only the inverse of the regularizer
    'solver': ['saga'],
    'max_iter': [10000]
}

print(param_grid)

# we save the train and validation scores
# the validation scores are necessary to select the best model
# it's optional to save the train scores, it can be used to identify high
train_score = np.zeros(len(ParameterGrid(param_grid)))
val_score = np.zeros(len(ParameterGrid(param_grid)))
models = []

# loop through all combinations of hyperparameter combos
for p in range(len(ParameterGrid(param_grid))):
    params = ParameterGrid(param_grid)[p]
    print(' ',params)
    clf = LogisticRegression(**params,random_state = 42*i) # initialize
    clf.fit(X_train_prep,y_train) # fit the model
    models.append(clf) # save it
    # calculate train and validation accuracy scores
    y_train_pred = clf.predict(X_train_prep)
    train_score[p] = accuracy_score(y_train,y_train_pred)
    y_val_pred = clf.predict(X_val_prep)
    val_score[p] = accuracy_score(y_val,y_val_pred)
    print(' ',train_score[p],val_score[p])

# print out model parameters that maximize validation accuracy
print('best model parameters:',ParameterGrid(param_grid)[np.argmax(val_score)])
print('corresponding validation score:',np.max(val_score))
# collect and save the best model
final_models.append(models[np.argmax(val_score)])
# calculate and save the test score
y_test_pred = final_models[-1].predict(X_test_prep)
test_scores[i] = accuracy_score(y_test,y_test_pred)
print('test score:',test_scores[i])

```

```

randoms state 1
{'penalty': ['l1'], 'C': array([1.e-03, 1.e-02, 1.e-01, 1.e+00, 1.e+01, 1.e+
02, 1.e+03]), 'solver': ['saga'], 'max_iter': [10000]}
  {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.001)}
    0.80303030303030303 0.8015970515970516
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.01)}
    0.8421375921375921 0.8444410319410319
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.1)}
    0.8487919737919738 0.8539619164619164
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(1.0)}
    0.8510954135954136 0.8539619164619164
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
0.0)}
    0.8511977886977887 0.8541154791154791
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
00.0)}
    0.8509930384930385 0.8541154791154791
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
000.0)}
    0.8509930384930385 0.8541154791154791
best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 1000
0, 'C': np.float64(10.0)}
corresponding validation score: 0.8541154791154791
test score: 0.8524489482573315
randoms state 2
{'penalty': ['l1'], 'C': array([1.e-03, 1.e-02, 1.e-01, 1.e+00, 1.e+01, 1.e+
02, 1.e+03]), 'solver': ['saga'], 'max_iter': [10000]}
  {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.001)}
    0.80246723996724 0.8023648648648649
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.01)}
    0.8422911547911548 0.8372235872235873
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.1)}
    0.8510954135954136 0.8487407862407862
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(1.0)}
    0.8530405405405406 0.8513513513513513
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
0.0)}
    0.8529381654381655 0.851044226044226
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
00.0)}
    0.8530917280917281 0.8508906633906634
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
000.0)}
    0.8530917280917281 0.8508906633906634
best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 1000
0, 'C': np.float64(1.0)}
corresponding validation score: 0.8513513513513513
test score: 0.8541378780899739

```



```

randoms state 3
{'penalty': ['l1'], 'C': array([1.e-03, 1.e-02, 1.e-01, 1.e+00, 1.e+01, 1.e+
02, 1.e+03]), 'solver': ['saga'], 'max_iter': [10000]}
  {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.001)}
    0.8028767403767404 0.7959152334152334
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.01)}
    0.8447993447993448 0.8390663390663391
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.1)}
    0.851044226044226 0.8481265356265356
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(1.0)}
    0.853552416052416 0.8493550368550369
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
0.0)}
    0.8538083538083538 0.8488943488943489
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
00.0)}
    0.854013104013104 0.8487407862407862
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
000.0)}
    0.854013104013104 0.8487407862407862
best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 1000
0, 'C': np.float64(1.0)}
corresponding validation score: 0.8493550368550369
test score: 0.8536772608628896
randoms state 4
{'penalty': ['l1'], 'C': array([1.e-03, 1.e-02, 1.e-01, 1.e+00, 1.e+01, 1.e+
02, 1.e+03]), 'solver': ['saga'], 'max_iter': [10000]}
  {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.001)}
    0.8069717444717445 0.7976044226044227
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.01)}
    0.8463349713349714 0.8389127764127764
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.1)}
    0.854013104013104 0.8465909090909091
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(1.0)}
    0.8545761670761671 0.8481265356265356
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
0.0)}
    0.8539619164619164 0.8468980343980343
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
00.0)}
    0.854013104013104 0.8468980343980343
    {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
000.0)}
    0.8539619164619164 0.8467444717444718
best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 1000
0, 'C': np.float64(1.0)}
corresponding validation score: 0.8481265356265356
test score: 0.8473821587594043

```

```

randoms state 5
{'penalty': ['l1'], 'C': array([1.e-03, 1.e-02, 1.e-01, 1.e+00, 1.e+01, 1.e+
02, 1.e+03]), 'solver': ['saga'], 'max_iter': [10000]}
  {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.001)}
    0.8003685503685504 0.8043611793611793
  {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.01)}
    0.8435196560196561 0.8453624078624079
  {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(0.1)}
    0.8495085995085995 0.856418918918919
  {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64
(1.0)}
    0.8517096642096642 0.8579545454545454
  {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
0.0)}
    0.851965601965602 0.8574938574938575
  {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
00.0)}
    0.8518632268632269 0.8570331695331695
  {'solver': 'saga', 'penalty': 'l1', 'max_iter': 10000, 'C': np.float64(1
000.0)}
    0.8518632268632269 0.8570331695331695
best model parameters: {'solver': 'saga', 'penalty': 'l1', 'max_iter': 1000
0, 'C': np.float64(1.0)}
corresponding validation score: 0.8579545454545454
test score: 0.8487640104406572

```

Things to look out for

- are the ranges of the hyperparameters wide enough?
 - if you are unsure, save the training scores and plot the train and val scores!
 - do you see underfitting? model performs poorly on both training and validation sets?
 - do you see overfitting? model performs very good on training but worse on validation?
 - if you don't see both, expand the range of the parameters and you'll likely find a better model
 - read the manual and make sure you understand what the hyperparameter does in the model
 - some parameters (like regularization parameters) should be evenly spaced in log because there is no upper bound
 - some parameters (like l1_ratio for the elastic net) should be linearly spaced because they have clear lower and upper bounds
 - **if the best hyperparameter is at the edge of your range, you definitely need to expand the range if you can**
- not every hyperparameter is equally important
 - some parameters have little to no impact on train and validation scores

- visualize the results if in doubt
- is the best validation score similar to the test score?
 - it's usual that the validation score is a bit better than the test score
 - but if the difference between the two scores is significant over multiple random states, something could be off

Quiz

By the end of this lecture, you will be able to

- perform basic hyperparameter tuning
- **apply GridSearchCV**
- describe why interpretability is important

Hyperparameter tuning with folds

- the steps are a bit different

```
In [4]: from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import make_pipeline

df = pd.read_csv('../data/adult_data.csv')

# let's separate the feature matrix X, and target variable y
y = df['gross-income'] # remember, we want to predict who earns more than 50k
X = df.loc[:, df.columns != 'gross-income'] # all other columns are features

ordinal_ftrs = ['education']
ordinal_cats = [[' Preschool', ' 1st-4th', ' 5th-6th', ' 7th-8th', ' 9th', ' 10th',
                ' Some-college', ' Assoc-voc', ' Assoc-acdm', ' Bachelors', ' Ma
onehot_ftrs = ['workclass', 'marital-status', 'occupation', 'relationship', 'rac
minmax_ftrs = ['age', 'hours-per-week']
std_ftrs = ['capital-gain', 'capital-loss']

# collect all the encoders
preprocessor = ColumnTransformer(
    transformers=[
        ('ord', OrdinalEncoder(categories = ordinal_cats), ordinal_ftrs),
        ('onehot', OneHotEncoder(sparse_output=False, handle_unknown='ignore'),
        ('minmax', MinMaxScaler(), minmax_ftrs),
        ('std', StandardScaler(), std_ftrs)])

# all the same up to this point
```

```
In [5]: # we will use GridSearchCV and the parameter names need to contain the ML al
# the parameters of some ML algorithms have the same name and this is how we

param_grid = {
```

```

        'logisticregression__penalty': ['l1'],
        'logisticregression__solver': ['saga'],
        'logisticregression__max_iter': [10000],
        'logisticregression__C': np.logspace(-3,3,7) # only the inverse
    }

nr_states = 3
test_scores = np.zeros(nr_states)
final_models = []

for i in range(nr_states):
    # first split to separate out the test set
    # we will use kfold on other
    X_other, X_test, y_other, y_test = train_test_split(X,y,test_size = 0.2,

    # splitter for other
    kf = KFold(n_splits=4,shuffle=True,random_state=42*i)

    # the classifier
    clf = LogisticRegression(**params,random_state = 42*i) # initialize the

    # let's put together a pipeline
    # the pipeline will fit_transform the training set (3 folds), and transform
    # then it will train the ML algorithm on the training set and evaluate it
    # it repeats this step automatically such that each fold will be an evaluation
    pipe = make_pipeline(preprocessor,clf)

    # use GridSearchCV
    # GridSearchCV loops through all parameter combinations and collects the
    grid = GridSearchCV(pipe, param_grid=param_grid,scoring = 'accuracy',
                        cv=kf, return_train_score = True, verbose=True)

    # this line fits the model on other and loops through the 4 different values
    grid.fit(X_other, y_other)
    # save results into a data frame. feel free to print it and inspect it
    results = pd.DataFrame(grid.cv_results_)
    #print(results)

    print('best model parameters:',grid.best_params_)
    print('validation score:',grid.best_score_) # this is the mean validation score
    # save the model
    final_models.append(grid)
    # calculate and save the test score
    y_test_pred = final_models[-1].predict(X_test)
    test_scores[i] = accuracy_score(y_test,y_test_pred)
    print('test score:',test_scores[i])

```

```

Fitting 4 folds for each of 7 candidates, totalling 28 fits
best model parameters: {'logisticregression__C': np.float64(100.0), 'logisticregression__max_iter': 10000, 'logisticregression__penalty': 'l1', 'logisticregression__solver': 'saga'}
validation score: 0.8529253685503686
test score: 0.8479963150621833
Fitting 4 folds for each of 7 candidates, totalling 28 fits
best model parameters: {'logisticregression__C': np.float64(1.0), 'logisticregression__max_iter': 10000, 'logisticregression__penalty': 'l1', 'logisticregression__solver': 'saga'}
validation score: 0.8498925061425061
test score: 0.8581298940580377
Fitting 4 folds for each of 7 candidates, totalling 28 fits
best model parameters: {'logisticregression__C': np.float64(1.0), 'logisticregression__max_iter': 10000, 'logisticregression__penalty': 'l1', 'logisticregression__solver': 'saga'}
validation score: 0.852311117936118
test score: 0.8506064793489944

```

In [6]: results

Out[6]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_logisticregre
0	0.141838	0.011799	0.009326	0.000240	
1	0.296977	0.028689	0.009311	0.000110	
2	1.522729	0.272064	0.009537	0.000155	
3	8.284712	0.586564	0.009746	0.000292	
4	16.500764	1.776865	0.009721	0.000058	
5	21.036766	2.211653	0.009714	0.000211	
6	21.543215	2.167504	0.009744	0.000130	

7 rows × 22 columns

Things to look out for

- less code but more stuff is going on in the background hidden from you
 - looping over multiple folds
 - .fit_transform and .transform is hidden from you
- nevertheless, GridSearchCV and pipelines are pretty powerful
- working with folds is a bit more robust because the best hyperparameter is selected based on the average score of multiple trained models

Quiz

Can we use GridSearchCV with sets prepared by train_test_split in advance? Use the sklearn manual or stackoverflow to answer the question.

By the end of this lecture, you will be able to

- perform basic hyperparameter tuning
- apply GridSearchCV
- **describe why interpretability is important**

Intro to interpretability

Example 1:

- A bank uses ML to review loans.

	Requested	Received	Interest Rate
Person A	10k	8k	4%
Person B	10k	10k	5%
Person C	10k	2k	5%

- Which person received the worst outcome?
- All three applicants have similar financial backgrounds
- Person C is a member of the protected class but persons A and B are not
- Person C sues the bank for discrimination
- How can the bank check whether the algorithm does not discriminate against protected classes of people?

Example 2:

- An ML model predicts that a patient is at risk of heart disease
- Doctors will ask why?
- The 'why' is very important here
- Should the prediction be trusted?
 - If the model predicts heart disease because of the patient's zip code, the clinician will push back
 - bias or some spurious correlation is suspected, model needs to be revised
 - If model predict high cholesterol is the main reason, that aligns with medical knowledge
- An interpretable outcome can inform next steps!

- If high cholesterol is the main reason for the predicted heart disease risk, diet, life style changes, and statin might be the solution
- If family history is the main driver of the prediction, earlier screenings might be the solution

Interpretable or explainable ML (XML or XAI)

There are two main types:

- global explanations
 - does the model make predictions based on reasonable features?
 - one value per feature, it is a vector of shape (n_{ftrs})
 - it describes how important each feature is generally
 - good start but cannot be used to explain predictions of one datapoint!
- local explanations
 - can we trust the model's prediction for one specific data point?
 - one value per feature and data point, it is a 2D array with a shape of $(n_{\text{points}}, n_{\text{ftrs}})$ - the same shape as your feature matrix
 - it describes how important each feature is for predicting one particular data point
 - required when working with human data!

Mud card